

ENGINE

Efficient Tuning and Inference for Large Language Models on Textual Graphs

Yun Zhu^{1,*}, Yaoke Wang^{1,*}, Haizhou Shi² and Siliang Tang^{1,†}

<https://arxiv.org/pdf/2401.15569>

Key Points


- Using LLM for enhanced textual encoding and integration with GNNs.
- parameter- and memory-efficient fine-tuning method
- Combining the LLMs and GNNs through a **tunable side structure**.
- caching and dynamic early exit
- Freezing weights of the LLM

Earlier methods (A) [Static Embeddings]

static shallow embedding

struggle to capture context-aware information and complex semantic relationships

(e.g., bag-of-words, skip-gram)



(a)

Figure 2: Zhu et al [1].


Earlier methods (B) [Cascading]

combine LMs and GNNs in a cascading

the initial step involves finetuning pre-trained language models on downstream tasks

enhances the meaningfulness of node embeddings by fine-tuning LMs using neighbor information.

initially fine-tuning LMs through graph-related tasks (e.g., node classification, link prediction)



(b)


Figure 2: Zhu et al [1].

Earlier methods (C) [Iterative]

combine LMs and GNNs in an iterative or co-training structure

these co-training paradigms face significant scalability issues.

Encoding all neighbors by LMs introduces high-cost fine-tuning and inference overhead due to the substantial number of parameters in language models



(c)

Figure 2: Zhu et al [1].

3.1 Notations

Given a textual graph $G = \{\mathcal{V}, \{t_n\}_{n \in \mathcal{V}}, A, Y\}$, where \mathcal{V} is the node set consisting of N instances, $t_n \in \mathcal{T}^{Q_n}$ represents a sequential text for its node $n \in \mathcal{V}$, \mathcal{T} is the tokens dictionary, and Q_n is the sequence length, $A \in \mathbb{R}^{N \times N}$ denotes the adjacency matrix, and Y denotes labels for each node. To enhance scalability, a sampling function $\Gamma(\cdot)$ is applied to a large graph to obtain a set of small subgraphs $\{G_n\}_{n \in \mathcal{V}}$, where G_n represents the subgraph for the target node $n \in \mathcal{V}$. In this study, the focus is on the node classification task of textual graphs. Specifically, given a set of training nodes \mathcal{V}_{tr} , a classification model is trained on these nodes and evaluated on the remaining test nodes \mathcal{V}_{te} . Formally, given a set of training nodes and their induced subgraphs $\mathcal{G}^{tr} = \{G_n\}_{n \in \mathcal{V}_{tr}}$, the optimal predictor f_{θ^*} is formulated as

$$f_{\theta^*} \in \arg \max_{\theta} \mathbb{E}_{G_n \in \mathcal{G}^{tr}} P_{\theta}(\hat{y}_n = y_n \mid G_n), \quad (1)$$

where y_n denotes the true label of target node $n \in \mathcal{V}_{tr}$ and \hat{y}_n is the predicted label. It is noteworthy that our method can be

$$H^l = \text{LLM_Layer}^l(H^{l-1}), \quad (2)$$

where $\text{LLM_Layer}^l(\cdot)$ denotes the l -th layer of LLM, and $H^l \in \mathbb{R}^{B \times Q \times D}$ means the token-level representations in i -th layer. In the first layer, the input H^0 is equivalent to \mathcal{B} .

To achieve this, a readout function \mathcal{R} , such as mean pooling [Mesquita *et al.*, 2020], is applied to token-level representations:

$$z_i^l = \mathcal{R}(h_{i,1}^l, h_{i,2}^l, \dots, h_{i,Q}^l), \quad (3)$$

where $z_i^l \in \mathbb{R}^{1 \times D}$ denotes the representation of node $i \in \mathcal{V}_{tr}$.

Then these node-level representations are fed into GNN Ladders (*G-Ladders*) to improve the quality of node embeddings through structural information, ensuring accurate comprehension of the node’s semantics from a global perspective. The improved node-level representations \hat{Z}^l are derived as

$$\hat{Z}^l = \lambda^l \cdot \text{GNN}^l(\mathcal{P}^l(Z^l), A) + (1 - \lambda^l) \cdot \hat{Z}^{l-1}, \quad (4)$$

where $\mathcal{P}^l(\cdot) : \mathbb{R}^{B \times D} \rightarrow \mathbb{R}^{B \times K}$ is a projector that maps the node-level representations into low dimensions ($K \ll D$), thereby reducing the subsequent

final performance (refer to Table 11 in Appendix). λ^l is a learnable coefficient for combining information in the current layer l and the previous layer $l - 1$. It is modeled by

$$\mathcal{L} = \mathbb{E}_{i \in \mathcal{V}_{tr}} \text{CE}(\hat{y}_i, y_i), \text{ where } \hat{y}_i = \mathcal{C}(\hat{z}_i^L).$$

Model Architecture

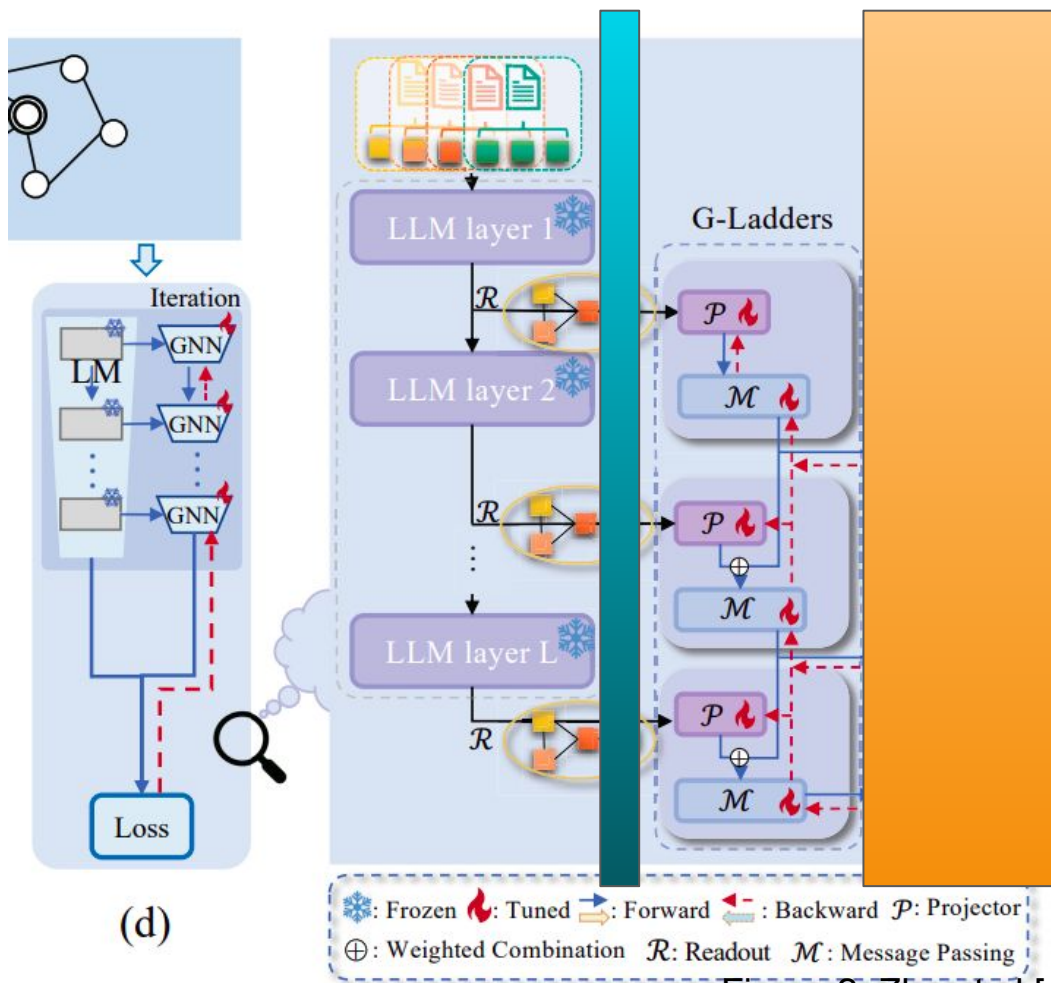


Figure 2: Zhu et al [1].

Caching

G-Ladders with LLMs via a side structure, eliminating the need for backpropagation through the LLMs.

allows us to precompute and cache node embeddings for reuse, further boosting computational efficiency

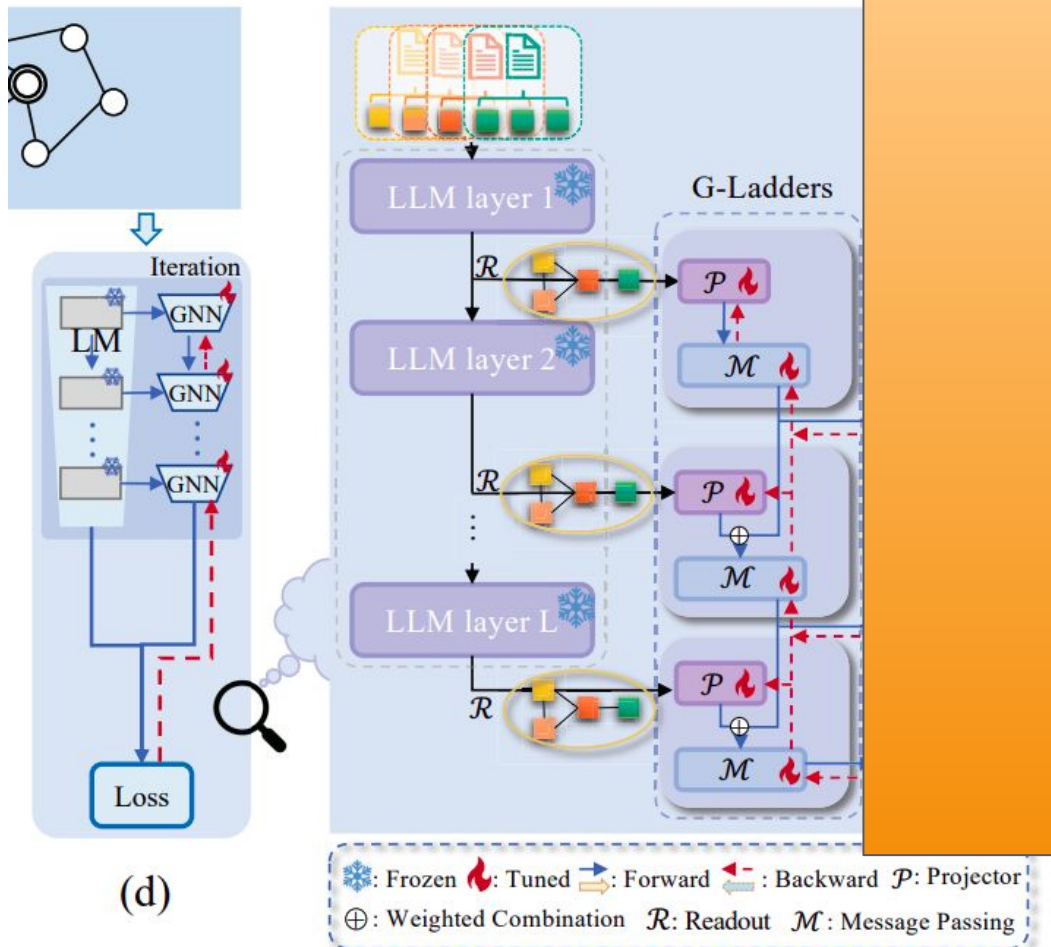


Figure 2: Zhu et al [1].

Early Stopping

with our method. Specifically, we add a lightweight single-layer MLP as an early exit classifier \mathcal{C}^l after each *G-Ladder*. During the model tuning, these classifiers are directly connected to the downstream task's training objective, *e.g.*, the cross-entropy loss between the true label y :

$$C^{l*} \in \arg \min_{\mathcal{C}^l} \mathbb{E}_{i \in \mathcal{V}_{tr}} \text{CE}(\mathcal{C}^l(\hat{z}_i^l), y_i). \quad (6)$$

Patience-based criteria imply that if consecutive p early exit classifiers predict the same results, where p represents the number of layers.

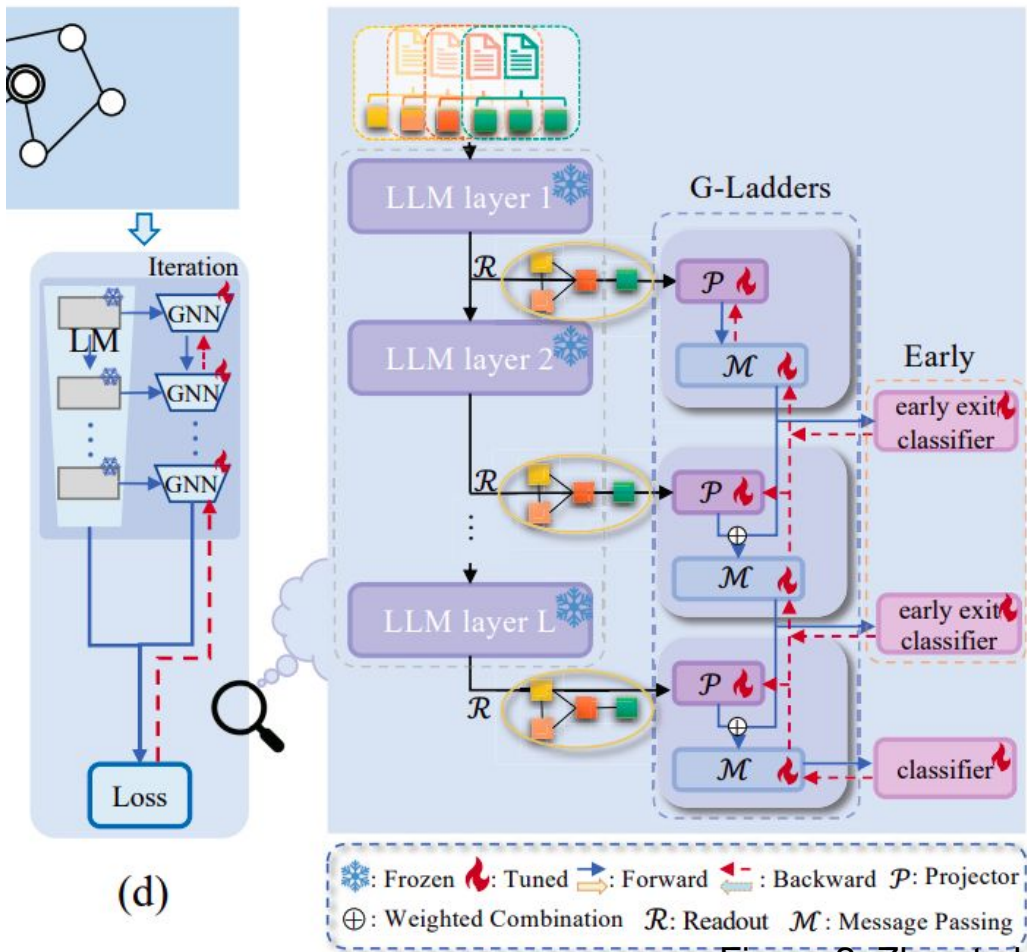


Figure 2: Zhu et al [1].

bibliography

1. Efficient Tuning and Inference for Large Language Models on Textual Graphs
(<https://arxiv.org/pdf/2401.15569>)